

Real-world Multi-object, Multi-grasp Detection

Fu-Jen Chu, Ruinian Xu and Patricio A. Vela

Abstract—A deep learning architecture is proposed to predict graspable locations for robotic manipulation. It considers situations where no, one, or multiple object(s) are seen. By defining the learning problem to be classification with null hypothesis competition instead of regression, the deep neural network with RGB-D image input predicts multiple grasp candidates for a single object or multiple objects, in a single shot. The method outperforms state-of-the-art approaches on the Cornell dataset with 96.0% and 96.1% accuracy on image-wise and object-wise splits, respectively. Evaluation on a multi-object dataset illustrates the generalization capability of the architecture. Grasping experiments achieve 96.0% grasp localization and 89.0% grasping success rates on a test set of household objects. The real-time process takes less than .25s from image to plan.

Index Terms—Perception for Grasping; Grasping; Deep Learning in Robotic Automation

I. INTRODUCTION

WHILE manipulating objects is relatively easy for humans, reliably grasping arbitrary objects remains an open challenge for robots. Resolving it would advance the application of robotics to industrial use cases, such as part assembly, binning, and sorting. Likewise, it would advance the area of assistive robotics, where the robot interacts with its surroundings in support of human needs. Robotic grasping involves perception, planning, and control. As a starting point, knowing which object to grab and how to do so are essential aspects. Consequently, accurate and diverse detection of robotic grasp candidates for target objects should lead to a better grasp path planning and improve the overall performance of grasp-based manipulation tasks.

The proposed solution utilizes a deep learning strategy for identifying suitable grasp configurations from an input image. In the past decade, deep learning has achieved major success on detection, classification, and regression tasks [1]–[3]. Its key strength is the ability to leverage large quantities of labelled and unlabelled data to learn powerful representations without hand-engineering the feature space. Deep neural networks have been shown to outperform hand-designed features and reach state-of-the-art performance.

In this research problem, we are interested in tackling the problem of identifying viable candidate robotic grasps of objects in a RGB-D image. The envisioned gripper is a parallel plate gripper (or similar in functionality). The principal difficulty comes from the variable shapes and poses



Fig. 1. Simultaneous multi-object, multi-grasp detection by the proposed model. Training used the Cornell dataset with the standard object-wise split. The red lines correspond to parallel plates of the grasping gripper. The white lines indicate the distance between the plates before the grasp is executed.

of objects as imaged by a camera. The structural features defining successful grasps for each object may be different; all possible features should be identified through the learning process. The proposed architecture associated to the grasp configuration estimation problem relies on the strengths of deep convolutional neural networks (CNNs) at detection and classification. Within this architecture, the identification of grasp configuration for objects is broken down into a grasp detection processes followed by a more refined grasp orientation classification process, both embedded within two coupled networks.

The proposed architecture includes a grasp region proposal network for identification of potential grasp regions. The network then partitions the grasp configuration estimation problem into regression over the bounding box parameters, and classification of the orientation angles, from RGB-D data. Importantly, the orientation classifier also includes a *No Orientation* competing class to reject spuriously identified regions for which no single orientation classification performs well, and to act as a competing no-grasp class. The proposed approach predicts grasp candidates in more realistic situations where no, single, and multiple objects may be visible; it also predicts multiple grasps with confidence scores (see Fig. 1, scores not shown for clarity). A new multi-objects grasp dataset is collected for evaluation with the traditional performance metric of false-positives-per-image. We show how the multi-grasp output and confidence scores can inform a subsequent planning process to take advantage of the multiple outputs for improved grasping success.

The main contributions of this paper are threefold:

- (1) A deep network architecture that predicts multiple grasp candidates in situations when none, single or multiple objects are in the view. Compared to baseline methods, the classification-based approach demonstrates improved out-

Manuscript received: February 24, 2018; Revised May 15, 2018; Accepted June 14, 2018.

This paper was recommended for publication by Editor Han Ding upon evaluation of the Associate Editor and Reviewers' comments. This work was supported in part by NSF Award #1605228.

Fu-Jen Chu, Ruinian Xu and Patricio A. Vela are with Institute for Robotics and Intelligent Machines, Georgia Institute of Technology, GA, USA. {fujenchu, gtsimonxu, pvela}@gatech.edu

Digital Object Identifier (DOI): see top of this page.

comes on the Cornell dataset [4] benchmark, achieving state-of-the-art performance on image-wise and object-wise splits.

(2) A multi-object, multi-grasp dataset is collected and manually annotated with grasp configuration ground-truth as the Cornell dataset. We demonstrate the generalization capabilities of the architecture and its prediction performance on the multi-grasp dataset with respect to false grasp candidates per image versus grasp miss rate. The dataset is available at github.com/ivalab/grasp_multiObject.

(3) Experiments with a 7 degree of freedom manipulator and a time-of-flight RGB-D sensor quantify the system's ability to grasp a variety of household objects placed at random locations and orientations. Comparison to published works shows that the approach is effective, achieving a sensible balance for real-time object pick-up with an 89% success rate and less than 0.25 s from image to prediction to plan.

II. RELATED WORK

Research on grasping has evolved significantly over the last two decades. Altogether, the review papers [5]–[8] provide a good context for the overall field. This section reviews grasping with an emphasis on learning-based approaches and on representation learning.

Early work on perception-based learning approaches to grasping goes back to [9], which identified a low dimensional feature space for identifying and ranking grasps. Importantly it showed that learning-based methods could generalize to novel objects. Since then, machine learning methods have evolved alongside grasping strategies. Exploiting the input/output learning properties of machine learning (ML) systems, [10] proposed to learn the image to grasp mapping through the manual design of convolutional networks. As an end-to-end system, reconstruction of the object's 3D geometry is not needed to arrive at a grasp hypothesis. The system was trained using synthetic imagery, then demonstrated successful grasping on real objects. Likewise, [11] employed a CNN-like feature space with random forests for grasp identification. In addition to where to grasp, several efforts learn the grasp approach or pre-grasp strategy [12], [13], while some focus on whether the hypothesized grasp is likely to succeed [14]. Many of these approaches exploited contemporary machine learning algorithms with manually defined feature spaces.

At the turn of this decade, two advances led to new methods for improving grasp identification: the introduction of low-cost depth cameras, and the advent of computational frameworks to facilitate the construction and training of CNNs. The advent of consumer depth cameras enabled models of grasping mapping to encode richer features. In particular [15] represented grasps as a 2D oriented rectangle in the image space with the local surface normal as the approaching vector; this grasp configuration vector has been adopted as the well-accepted formulation. Generally, the early methods using depth cameras sought to recover the 3D geometry from point clouds for grasp planning [16], with manually derived feature space used in the learning process. Deep learning approaches arrived later [2] and were quickly adopted by the computer vision community [17].

Deep learning avoids the need for engineering feature spaces, with the trade-off that larger datasets are needed. The trade-off is usually mitigated through the use of pre-training on pre-existing computer vision datasets followed by fine-tuning on a smaller, problem-specific dataset. Following a sliding window approach, [18] trained a two stage multi-modal network, with the first stage generating hypotheses for a more accurate second stage. Similarly, [19] first performed an image-wide pre-processing step to identify candidate object regions, followed by application of a CNN classifier for each region. To avoid sliding windows or image-wide search, end-to-end approaches are trained to output a single grasp configuration from the input data [20]–[22]. Regression-based approaches [20] require compensation through image partitioning since the grasp configuration space is non-convex. Following the architecture in [20], [22] experimented on real objects with physical grasping experiments. The two-stage network in [21] first output a learnt feature, which was then used to provide a single grasp output. These approaches may suffer from averaging effects associated to the single-output nature of the mapping. Guo et al. [23] employed a two-stage process with a feature learning CNN, followed by specific deep network branches (graspable, bounding box, and orientation). Alternatively, [24], [25] predicted grasp scores over the set of all possible grasps (which may be discretized). Such networks admit the inclusion of end-effector uncertainty. Most deep network approaches mentioned above start with the strong prior that every image contains a single object with a single grasp target (except [19]). This assumption does not generically hold in practice: many objects have multiple grasp options and a scene may contain more than one object.

Another line of research is to learn the mapping from vision input to robot motion to achieve grasping. To directly plan grasps, Lu et al. [26] proposed to predict and maximize grasp success by inferring grasp configurations from vision input for grasp planning. Research on empirical grasp planning with reinforcement learning (RL) acquired samples from robots in real experiments [27]. The training time involved several weeks and led to limitation of its scalability. The work [28] collected over 800k data points with up to 14 robotic arms running in parallel for learning visual servoing. The training time involved over 2 months. Generalization performance of RL solution to environmental changes remains unknown.

Inspired by [29], we propose to incorporate a *grasp region proposal network* to generate candidate regions for feature extraction. Furthermore, we propose to transform grasp configuration from a regression problem formulated in previous works [20], [21] into a combination of region detection and orientation classification problems (with null hypothesis competition). We utilize ResNet [17], the current state-of-the-art deep convolutional neural network, for feature extraction and grasp prediction. Compared to previous approaches, our method considers more realistic scenarios with multiple objects in a scene. The proposed architecture predicts multiple grasps with corresponding confidence scores, which aids the subsequent planning process and actual grasping.




Fig. 2. (a) The 5D grasp representation. (b) A grasp rectangle is first set to the zero orientation for grasp proposal training. The angle θ is one of the discrete rotation angles. (c) Each element in the feature map is an anchor and corresponds to multiple candidate grasp proposal bounding boxes.

III. PROBLEM STATEMENT

Given corresponding RGB and depth images of a novel object, the objective is to identify the grasp configurations for potential grasp candidates of an object for the purpose of manipulation. The 5-dimensional grasp rectangle is the grasp representation employed [15]. It is a simplification of the 7-dimensional representation [15] and describes the location, orientation, and opening distance of a parallel plate gripper prior to closing on an object. The 2D orientated rectangle, shown in Fig. 2a depicts the gripper's location (x, y) , orientation θ , and opening distance h . An additional parameter describing the length w completes the bounding box grasp configuration,

$$g = \{x, y, \theta, w, h\}^T. \quad (1)$$

Thinking of the center of the bounding box with its local (x, y) axes aligned to the w and h variables, respectively, the first three parameters represent the $SE(2)$ frame of the bounding box in the image, while the last two describe the dimensions of the box.

IV. APPROACH

Much like [20], [23], the proposed approach should avoid sliding window such as in [18] for real-time implementation purposes. We avoid the time-consuming sliding-window approach by harnessing the capacity of neural networks to perform bounding box regression, and thereby to predict candidate regions on the full image directly. Furthermore, we preserve all possible grasps and output all ranked candidates, instead of regressing a single outcome. To induce a richer feature representation and learn more of the structural cues, we propose to use a deeper network model compared to previous works [18], [20], [27], with the aim of improving feature extraction for robotic grasp detection. We adopt the ResNet-50 [17] with 50 layers, which has more capacity and should learn better than the AlexNet [2] used in previous works (8 layers). ResNet is known for its residual learning concept to overcome the challenge of learning mapping functions. A residual block is designed as an incorporation of a skip connection with standard convolutional neural network. This design allows the block to bypass the input, and encourage convolutional layers to predict the residual for the final mapping function of a residual block.

The next three subsections describe the overall architecture of the system. It includes integration of the proposal network with a candidate grasp region generator; a description of our choice to define grasp parameter estimation as a combination of regression and classification problems; and an explanation of the multi-grasp detection architecture.

A. Grasp Proposals

The first stage of the deep network aims to generate grasp proposals across the whole image, avoiding the need for a separate object segmentation pipeline.

Inspired by *Region Proposal Network* (RPN) [29], the *Grasp Proposal Network* in our architecture (Fig. 3) works as RPN and shares a common feature map ($14 \times 14 \times 1024$ feature map) of intermediate convolutional layers from ResNet-50 (layer 40). The *Grasp Proposal Network* outputs a $1 \times 1 \times 512$ feature which is then fed into two sibling fully connected layers. The two outputs specify both probability of grasp proposal and proposal bounding box for each of r anchors on the shared feature map. The ROI layer extracts features with corresponding proposal bounding boxes and sends to the rest of the networks.

The *Grasp Proposal Network* works as sliding a mini-network over the feature map. At each anchor of the feature map, by default 3 scales and 3 aspect ratios are used for grasp reset bounding box shape variations, as shown in Fig 2c. Hence $r \times 3 \times 3$ predictions would be generated in total. For ground truth, we reset each orientated ground truth bounding box to have vertical height and horizontal width, as shown in Fig. 2b. Let t_i denote the 4-dimensional vector specifying the reset (x, y, w, h) of the i -th grasp configuration, and p_i denote the probability of the i -th grasp proposal. For the index set of all proposals \mathbf{I} , we define the loss of grasp proposal net (gpn) to be:

$$L_{gpn}(\{(p_i, t_i)_{i=1}^{\mathbf{I}}\}) = \sum_i L_{gp_cls}(p_i, p_i^*) + \lambda \sum_i p_i^* L_{gp_reg}(t_i, t_i^*). \quad (2)$$

where L_{gp_cls} is the cross entropy loss of grasp proposal classification (gp_cls), L_{gp_reg} is the l_1 regression loss of grasp proposal (gp_reg) with weight λ . We denote $p_i^* = 0$ for no grasp and $p_i^* = 1$ when a grasp is specified. The variable t_i^* is the ground truth grasp coordinate corresponding to p_i^* .

Compared to the widely applied selective search used in R-CNN [30], RPN learns object proposals end-to-end from the input without generating region of interests beforehand. This latter, streamlined approach is more applicable to real-time robotic applications.

B. Grasp Orientation as Classification

Many prior approaches [20], [21] regress to a single 5-dimensional grasp representation $g = \{x, y, w, h, \theta\}$ for a RGB-D input image. Yet to predict either on $SE(2)$ (planar pose) or on S^1 (orientation) involves predicting coordinates that lies in a non-Euclidean (non-convex) space where regression and its standard L2 loss may not perform well.




Fig. 3. Complete structure of our multi-object multi-grasp predictor. The network takes RG-D inputs, and predicts multiple grasps candidates with orientations and rectangle bounding boxes for each object in the view. Blue blocks indicate network layers and gray blocks indicate images and feature maps. Green blocks show the two loss functions. The grasp proposal network slides across anchors of intermediate feature maps from ResNet-50 with $k = 3 \times 3$ candidates predicted per anchor. The black lines of output bounding boxes denote the open length of a two-fingered gripper, while the red lines denote the parallel plates of the gripper.

Rather than performing regression, our multi-grasp localization pipeline quantizes the grasp representation orientation coordinate θ into R equal-length intervals (each interval is represented by its centroid), and formulates the input/output mapping as a classification task for grasp orientation. It differs from [23] in that we add a non-grasp collecting orientation class for explicit competition with a null hypothesis. If none of the orientation classifiers outputs a score higher than the non-grasp class, then the grasp proposal is considered incorrect and rejected. In contrast, [23] has a separate grasp confidence score, which may not capture well the orientation-dependent properties of grasps. The value of the non-grasp class is that it is necessary for the downstream multi-object, multi-grasp component of the final algorithm. The total number of classes is $|C| = R + 1$. Denote by $\{(l_i, \theta_i)\}_{i=1}^I$ where the i -th grasp configuration with classification label $l_i \in 1, \dots, R$ is associated with the angle θ_i . For the case of no possible orientation (i.e., the region is not graspable), the output label is $l = 0$ and there is no associated orientation. In this paper, $R = 19$ is utilized.

C. Multi-Grasp Detection

After the region proposal stage of the deep network, the last stage identifies candidate grasp configurations. This last stage classifies the predicted region proposals from previous stage into R regions for grasp configuration parameter θ . At the same time the last stage also refines the proposal bounding box to a non-oriented grasp bounding box (x, y, w, h) .

To process the region proposals efficiently, we integrate an *ROI pooling layer* [31] into ResNet-50 so that it may share ResNet's convolutional layers. Sharing the feature map with previous layers avoids re-computation of features within the region of interest. An *ROI pooling layer* stacks all of the features of the identified grasp proposals, which then get fed to two sibling fully connected layers for orientation parameter classification l and bounding box regression (x, y, w, h) . The ROI pooling layer receives its input from the intermediate convolutional layer of ResNet-50 (layer 40).

Let p_l denote the probability of class l after a softmax layer, and β_l denote the corresponding predicted grasp bounding box. Define the loss function of the grasp configuration prediction

(gcr) to be:

$$L_{gcr}(\{(p_l, \beta_l)\}_{c=0}^C) = \sum_{c=0}^C L_{gcr_cls}(p_l) + \lambda_2 \sum_{c=0}^C \mathbf{1}_{c \neq 0}(c) L_{gcr_reg}(\beta_c, \beta_c^*). \quad (3)$$

where L_{gcr_cls} is the cross entropy loss of grasp angle classification (gcr_cls), L_{gcr_reg} is the l_1 regression loss of grasp bounding boxes (gcr_reg) with weight λ_2 , and β_c^* is the ground truth grasp bounding box.

With the modified ResNet-50 model, end-to-end training for grasp detection and grasp parameter estimation employs the total loss:

$$L_{total} = L_{gpn} + L_{gcr}. \quad (4)$$

The streamlined system generates grasp proposals at the ROI layer, stacks all ROIs using the shared feature, and the additional neurons of the two sibling layers output grasp bounding boxes and orientations, or reject the proposal.

V. EXPERIMENTS AND EVALUATION

Evaluation of the grasp identification algorithm utilizes the Cornell Dataset for benchmarking against other state-of-the-art algorithms. To demonstrate the multi-object, multi-grasp capabilities, a new dataset is carefully collected and manually annotated. Both datasets consist of color and depth images for multiple modalities. In practice, not all possible grasps are covered by the labelled ground truth, yet the grasp rectangles are comprehensive and representative for diverse examples of good candidates. The scoring criteria takes into account the potential sparsity of the grasp configuration by including an acceptable proximity radius to the ground truth grasp configuration.

a) *Cornell Dataset*: The Cornell Dataset [4] consists of 885 images of 244 different objects, with several images taken of each object in various orientations or poses. Each distinct image is labelled with multiple ground truth grasps corresponding to possible ways to grab the object.

b) *Multi-Object Dataset*: Since the Cornell Dataset scenarios consist of one object in one image, we collect a Multi-Object Dataset for the evaluation of the multi-object/multi-grasp case. Our dataset is meant for evaluation and consists of 96 images with 3-5 different objects in a single image. We follow the same protocol as the Cornell Dataset by taking several images of each set of objects in various orientations or poses. Multiple ground truth grasps for each object in each image are annotated using the same configuration definition.

A. Cornell Data Preprocessing

To reuse the pre-trained weights of ResNet-50 on COCO-2014 dataset [32], the Cornell dataset is preprocessed to fit the input format of the ResNet-50 network. For comparison purposes, we follow the same procedure in [20] and substitute the blue channel with the depth channel. Since RGB data lies between 0 to 255, the depth information is normalized to the same range. The mean image is chosen to be 144, while the pixels on the depth image with no information were replaced with zeros. For the data preparation, we perform extensive data augmentation. First, the images are center cropped to obtain a 351x351 region. Then the cropped image is randomly rotated between 0 to 360 degree and center cropped to 321x321 in size. The rotated image is randomly translated in x and y direction by up to 50 pixels. The preprocessing generates 1000 augmented data for each image. Finally the image is resized to 227x227 to fit the input of ResNet-50 architecture.

B. Pre-Training

To avoid over-fitting and precondition the learning process, we start with the pretrained ResNet-50. As shown in Fig 3, we implement grasp proposal layer after the third residual block by sharing the feature map. The proposals are then sent to the ROI layer and fed into the fourth residual layer. The 7×7 average pool outputs are then fed to two fully connected layers for final classification and regression. All new layers beyond ResNet-50 are trained from scratch.

Because the orientation is specified as a class label and assigned a specific orientation, the Cornell dataset needs to be converted into the expected output format of the proposed network. We equally divide 180 degrees into R regions (due to symmetry of the gripper) and assign the continuous ground truth orientation to the nearest discrete orientation.

C. Training

For training, we train the whole network end-to-end for 5 epochs on a single nVidia Titan-X (Maxwell architecture). The initial learning rate is set to 0.0001. And we divide the learning rate by 10 every 10000 iterations. Tensorflow is the implementation framework with cudnn-5.1.10 and cuda-8.0 packages. The code will be publicly released.

D. Evaluation Metric

Accuracy evaluation of the grasp parameters involves checking for proximity to the ground truth according to established criteria [20]. A candidate grasp configuration is reported correct if both:

- 1) the difference of angle between predicted grasp g_p and ground truth g_t is within 30° , and
- 2) the Jaccard index of the predicted grasp g_p and the ground truth g_t is greater than 0.25, e.g.,

$$J(g_p, g_t) = \frac{|g_p \cap g_t|}{|g_p \cup g_t|} > 0.25 \quad (5)$$

The Jaccard index is similar to the Intersection over Union (IoU) threshold for object detection.

TABLE I
SINGLE-OBJECT SINGLE-GRASP EVALUATION

approach	image-wise	object-wise	speed
	Prediction Accuracy (%)		fps
Jiang et al. [15]	60.5	58.3	0.02
Lenz et al. [18]	73.9	75.6	0.07
Redmon et al. [20]	88.0	87.1	3.31
Wang et al. [19]	81.8	N/A	7.10
Asif et al. [11]	88.2	87.5	–
Kumra et al. [21]	89.2	88.9	16.03
Mahler et al. [25]	93.0	N/A	~1.25
Guo et al. [23]	93.2	89.1	–
Ours: VGG-16 (RGB-D)	95.5	91.7	17.24
Ours: Res-50 (RGB)	94.4	95.5	8.33
Ours: Res-50 (RGB-D)	96.0	96.1	8.33

TABLE II
PREDICTION ACCURACY (%) AT DIFFERENT JACCARD THRESHOLDS

split	0.25	0.30	0.35	0.40
image-wise	96.0	94.9	92.1	84.7
object-wise	96.1	92.7	87.6	82.6

VI. RESULTS

A. Single-object Single-grasp

Testing of the proposed architecture on the Cornell Dataset, and comparison with prior works lead to Table I. For this single-object/single-grasp test, the highest output score of all grasp candidates output is chosen as the final output. The proposed architecture outperforms all competitive methods. On image-wise split, our architecture reaches 96.0% accuracy; on object-wise split for unseen objects, 96.1% accuracy is achieved. We also tested our proposed architecture by replacing ResNet-50 with VGG-16 architecture, a smaller deep net with 16 layers. With VGG-16, our model still outperforms competitive approaches. Yet the deeper ResNet-50 achieve 4.4% more on unseen objects. Furthermore, we experiment on RGB images without depth information with ResNet-50 version and both image-wise and object-wise split perform slightly worse than our proposed approach, indicating the effectiveness of depth. The third column contains the run-time of methods that have reported it, as well as the runtime of the proposed method. Computationally, our architecture detects and localize multiple grasps in 0.120s, which is around 8 fps and is close to usable in real time applications. The VGG-16 architecture doubles the speed with some prediction accuracy loss.

Table II contains the outcomes of stricter Jaccard indexes for the ResNet-50 model. Performance decreases with stricter conditions but maintains competitiveness even at 0.40 IoU condition. Typical output of the system is given in Fig. 4a, where four grasps are identified. Limiting the output to a single grasp leads to the outputs depicted in Fig. 4b. In the multi-grasp case, our system not only predicts universal grasps learned from ground truth, but also contains candidate grasps not contained in the ground truth, Fig. 4c.

B. Single-object Multi-grasp

For realistic robotic application, a viable grasp usually depends both on the object and its surroundings. Given that




Fig. 4. Output 5D grasp configuration of system for Cornell dataset inputs: (a) the multiple grasp options output for an object; (b) the top grasp outputs for several objects; (c) output grasps (red) and ground-truth grasps (green) showing that the system may output grasps for which there is no ground truth; (d) multi-grasp output for several objects. The green rectangles are ground truth and the red rectangles represent predicted grasps for each unseen object.

one grasp candidate may be impossible to achieve, there is benefit to provide a rank ordered list of grasp candidates. Our system provides a list of high quality grasp candidates for a subsequent planner to select from. Fig. 4d shows samples of the predicted grasps and corresponding ground truths. To evaluate the performance of the multi-grasp detector, we employ the same scoring system as with the single grasp, then generate the miss rate as a function of the number of false positives per image (FPPI) by varying the detection threshold (see Fig. 5a for the single-object multi-grasp case). The model achieves 28% and 25% miss rate at 1 FPPI for object-wise split and image-wise split, respectively. A false positive means an incorrect grasp candidate for the object. Thus, accepting that there may be 1 incorrect candidate grasp per image, the system successfully detects 72% (75%) of possible grasps for object-wise (image-wise) split. The model performs slightly better in image-wise split than object-wise split due to unseen objects in the latter.

C. Multi-object Multi-grasp

Here, we apply the proposed architecture to a multi-object multi-grasp task using our Multi-Object dataset. The trained network is the same trained network we've been reporting the results for (trained only on the Cornell dataset with both image-split and object-split variants). Testing involves evaluating against the multi-object dataset, and represents a cross domain application with unseen objects.

Fig. 5a depicts the plot of miss rate versus FPPI. At 1FPPI, the system achieves 53% and 49% prediction accuracy with image-split model and object-split networks, respectively. Visualizations of predicted grasp candidates are depicted in Fig. 5b. The model successfully locates multiple grasp candidates on multiple new objects in the scene with very few false positives, and hence is practical for robotic manipulations.

D. Physical Grasping

To confirm and test the grasp prediction ability in practice, a physical grasping system is set up for experiments (see Fig. 5c). As in [22], performance is given for both the vision sub-system and the subsequently executed grasp movement. The dual scores aid in understanding sources of error for the overall experiment. To evaluate the vision sub-system, each RGB-D input of vision sub-system is saved to disk and annotated with the same protocol as the Cornell and Multi-Object

TABLE III
PHYSICAL GRASPING COMPARISON

approach	Top-1		Nearest to center	
	detected	physical	detected	physical
banana	10/10	7/10	10/10	8/10
glasses	9/10	8/10	9/10	9/10
ball	10/10	9/10	10/10	9/10
tape	10/10	10/10	10/10	10/10
screwdriver	9/10	7/10	9/10	7/10
stapler	10/10	10/10	10/10	9/10
spoon	9/10	9/10	10/10	10/10
bowl	10/10	10/10	10/10	10/10
scissors	9/10	8/10	9/10	9/10
mouse	9/10	8/10	9/10	8/10
average (%)	95.0	86.0	96.0	89.0

TABLE IV
PHYSICAL GRASPING EVALUATION ON SAME ROBOT AND OBJECTS

approach	Cornell splits	Physical grasp	
	image / object	detected	physical
Kumra et al. [21]	88.7 / 86.5	61/100	56/100
Guo et al. [23]	93.8 / 89.9	89/100	81/100
Ours (Top-1)	96.0 / 96.1	95/100	86/100
Ours (center)	96.0 / 96.1	96/100	89/100

datasets. The evaluation metric uses Jaccard index 0.25 and angle difference 30° thresholds. A set of 10 commonly seen objects was collected from Cornell dataset for the experiment. For each experiment, an object was randomly placed on a reachable surface at different locations and orientations. Each object was tested 10 times. The outcome of physical grasping was marked as pass or fail.

Table III shows the performance of both the vision sub-system and the physical grasping sub-system for two different policies. For the first (Top-1), we used the the grasp candidate with the highest confidence score. For the second, the planner chose the grasp candidate closest to the image-based center of the object from the top- N candidates ($N = 25$ in this experiment). In real-world physical grasping, grasp candidates close to the image-based centroid of object should be helpful, by creating a more balanced grasp for many objects. The lowest performing objects are those with a less even distribution of mass or shape (screwdriver), meaning that object-specific grasp priors coupled to the multi-grasp output might improve grasping performance. We leave this for future work as our system does not perform object recognition. Also, the tested




Fig. 5. Detection results of the system: (a) The ROC curves of our system on single-object multi-grasp scenario and multi-object multi-grasp scenario, respectively. The model was trained on Cornell Dataset and tested on our own multi-object dataset. (b) Detection results of our system on multi-object multi-grasp scenario. The model was trained on Cornell Dataset and tested on our own multi-object dataset. Red rectangle represents the predicted grasp on each unseen object. (c) Experiment setting for physical grasping test. The manipulator is a 7 degree of freedom redundant robotic arm. The vision device is META-1 AR glasses with time-of-flight for RGB-D input.

TABLE V
PHYSICAL GRASPING COMPARISON

approach	Time (s)		Settings		Success (%)
	detect	plan	object	trial	
[18]	13.50	–	30	100	84 / 89*
[22]	1.80	–	10	–	62
[27]	–	–	15	150	66
[26]	–	2~3	10	–	84
[25]	0.80	–	10	50	80
[25]+refits	2.50	–	40	100	94
Ours	0.12	0.10	10	100	89.0

* Outcomes are for Baxter / PR2 robots, respectively, with the difference arising from the different gripper spans.

objects are unseen.

For direct physical grasping comparisons, state-of-the-art approaches were implemented and applied to the same objects with the same robot. The reference approaches selected are [23] due to its performance on the Cornell dataset, and [21] due to the closely related back-bone architecture in the model design. Table IV reports both the accuracies of our implementations on standard Cornell dataset and physical grasping success. The set of objects is the same as in table III. The method in [21] has two parallel ResNet-50 for RGB and depth input, respectively. Our implementation achieves similar results on Cornell dataset. However, it overfits to the Cornell dataset as performance drops substantially for real-world objects (56% success rate). Our implementation of [23] achieves slightly better than reported, and reaches 81% success rate on physical grasping. Our proposed approach outperforms both across the board.

Table V further compares our experimental outcomes with state-of-the-art published works with physical grasping. The testing sets for reported experiments may include different object class/instance. Even though the object classes may be the same, the actual objects used could differ. Nevertheless, the comparison should provide some context for grasping performance and computational costs relative to other published approaches. The experiments in [22] had 6 objects in common with ours. On the common subset, [22] reports a 55.0% success rate with a 60s execution time, while ours achieves 86.7% with a 15s execution time (mostly a consequence of joint rate limits). The approach described in [25] reported

80.0% success rate on 10 household objects and 94.0% when using a cross entropy method [28] to sample and re-fit grasp candidates (at the cost of greater grasp detection time). The RL approach taking several weeks achieved 66.0% on seen and unseen objects [27]. Not included in the table are the reported results of [28], due to different experimental conditions. They reported 90% success on grasping objects from a bin with replacement, and 80% without replacement (100 trials using unseen objects). Our approach achieves 89.0% in real-time, with subsequent planning of the redundant manipulator taking 0.1 secs. Overall it exhibits a good balance between accuracy and speed for real world object grasping tasks.

E. Ablation Study

This section reviews a set of experiments, summarized in Table VI, examining the contributions of the proposed architecture's components. Firstly, ResNet-50 was used to regress RGB input to 5D grasp configuration output (a). This architecture can be recognized as [20] with a deeper network and without depth information. Then two ResNet-50 networks (b) processed RGB and depth data, respectively, with a small network regressing the concatenated feature for the grasp configuration. This architecture matches [21] and boosts performance. However, the doubled number of parameters results in difficulties when deploying on real-world grasping. To keep the architecture size, one color channel (blue) is replaced with depth information, while the performance is maintained (c). Next, grasp orientation is quantized and an extra branch is trained to classify grasping orientation of an object (d). The last two instances integrate grasp proposals into the ResNet-50 back-bone with added layers, for color (e) and RGD (f) input data. The multi-grasp outputs overcome averaging effects [20] without the need to separate an image into grids. The ablation study identifies the contribution of classification, grasp proposal and the selection policy. In addition, the RGB-only version of the proposed method is still able to achieve good performance, being slightly worse than including depth information.

VII. CONCLUSION

We presented a novel grasping detection system to predict grasp candidates for novel objects in RGB-D images. Com-

TABLE VI
ABLATION STUDY

Architecture	Cornell Splits		Number of Parameters
	image	object	
(a) RGB	86.4	85.4	24559685
(b) RGB + depth	88.7	86.5	51738757
(c) RGD	88.1	86.0	24559685
(d) RGD + cls*	89.8	89.3	24568919
(e) RGB + cls + gp	94.4	95.5	28184211
(f) RGD + cls + gp	96.0	96.1	28184211

* cls: classification; gp: grasp proposal

pared to previous works, our architecture is able to predict multiple candidate grasps instead of single outcome, which shows promise to aid a subsequent grasp planning process. Our regression as classification approach transforms orientation regression to a classification task, which takes advantage of the high classification performance of CNNs for improved grasp detection outcomes. We evaluated our system on the Cornell grasping dataset for comparison with state-of-the-art system using a common performance metric and methodology to show the effectiveness of our design. We also performed experiments on self-collected multi-object dataset for multi-object multi-grasp scenario. Acceptable grasp detection rates are achieved for the case of 1 false grasp per image. Physical grasping experiments show a small performance loss (8.3%) when physically grasping the object based on correct candidate grasps found. The outcomes might be improved by fusing the multi-grasp output with object-specific grasp priors, which we leave to future work. All code and data will be publicly released.

REFERENCES

- [1] L. Bo, X. Ren, and D. Fox, "Unsupervised feature learning for RGB-D based object recognition," in *Experimental Robotics*, 2013, pp. 387–402.
- [2] A. Krizhevsky, I. Sutskever, and G. E. Hinton, "Imagenet classification with deep convolutional neural networks," in *Advances in Neural Information Processing Systems*, 2012, pp. 1097–1105.
- [3] J. Ngiam, A. Khosla, M. Kim, J. Nam, H. Lee, and A. Y. Ng, "Multimodal deep learning," in *Proceedings of International Conference on Machine Learning*, 2011, pp. 689–696.
- [4] R. L. Lab, "Cornell grasping dataset," http://pr.cs.cornell.edu/grasping/rect_data/data.php, 2013, accessed: 2017-09-01.
- [5] K. B. Shimoga, "Robot grasp synthesis algorithms: A survey," *The International Journal of Robotics Research*, vol. 15, no. 3, pp. 230–266, 1996.
- [6] A. Bicchi and V. Kumar, "Robotic grasping and contact: A review," in *Proceedings of IEEE International Conference on Robotics and Automation*, 2000, pp. 348–353.
- [7] A. Sahbani, S. El-Khoury, and P. Bidaud, "An overview of 3D object grasp synthesis algorithms," *Robotics and Autonomous Systems*, vol. 60, no. 3, pp. 326–336, 2012.
- [8] J. Bohg, A. Morales, T. Asfour, and D. Kragic, "Data-driven grasp synthesis: A survey," *IEEE Transactions on Robotics*, vol. 30, no. 2, pp. 289–309, 2014.
- [9] I. Kamon, T. Flash, and S. Edelman, "Learning to grasp using visual information," in *Proceedings of IEEE International Conference on Robotics and Automation*, vol. 3, 1996, pp. 2470–2476.
- [10] A. Saxena, J. Driemeyer, and A. Y. Ng, "Robotic grasping of novel objects using vision," *The International Journal of Robotics Research*, vol. 27, no. 2, pp. 157–173, 2008.
- [11] U. Asif, M. Bennamoun, and F. A. Sohel, "RGB-D object recognition and grasp detection using hierarchical cascaded forests," *IEEE Transactions on Robotics*, 2017.
- [12] S. Ekvall and D. Kragic, "Learning and evaluation of the approach vector for automatic grasp generation and planning," in *Proceedings of IEEE International Conference on Robotics and Automation*, 2007, pp. 4715–4720.
- [13] K. Huebner and D. Kragic, "Selection of robot pre-grasps using box-based shape approximation," in *IEEE/RSJ International Conference on Intelligent Robots and Systems*, 2008, pp. 1765–1770.
- [14] Q. V. Le, D. Kamm, A. F. Kara, and A. Y. Ng, "Learning to grasp objects with multiple contact points," in *Proceedings of IEEE International Conference on Robotics and Automation*, 2010, pp. 5062–5069.
- [15] Y. Jiang, S. Moseson, and A. Saxena, "Efficient grasping from RGBD images: Learning using a new rectangle representation," in *Proceedings of IEEE International Conference on Robotics and Automation*, 2011, pp. 3304–3311.
- [16] D. Rao, Q. V. Le, T. Phoka, M. Quigley, A. Sudsang, and A. Y. Ng, "Grasping novel objects with depth segmentation," in *IEEE/RSJ International Conference on Intelligent Robots and Systems*, 2010, pp. 2578–2585.
- [17] K. He, X. Zhang, S. Ren, and J. Sun, "Deep residual learning for image recognition," in *IEEE conference on Computer Vision and Pattern Recognition*, 2016, pp. 770–778.
- [18] I. Lenz, H. Lee, and A. Saxena, "Deep learning for detecting robotic grasps," *The International Journal of Robotics Research*, vol. 34, no. 4–5, pp. 705–724, 2015.
- [19] Z. Wang, Z. Li, B. Wang, and H. Liu, "Robot grasp detection using multimodal deep convolutional neural networks," *Advances in Mechanical Engineering*, vol. 8, no. 9, p. 1687814016668077, 2016.
- [20] J. Redmon and A. Angelova, "Real-time grasp detection using convolutional neural networks," in *Proceedings of IEEE International Conference on Robotics and Automation*, 2015, pp. 1316–1322.
- [21] S. Kumra and C. Kanan, "Robotic grasp detection using deep convolutional neural networks," in *IEEE/RSJ International Conference on Intelligent Robots and Systems*, 2017.
- [22] J. Watson, J. Hughes, and F. Iida, "Real-world, real-time robotic grasping with convolutional neural networks," in *Conference Towards Autonomous Robotic Systems*, 2017, pp. 617–626.
- [23] D. Guo, F. Sun, H. Liu, T. Kong, B. Fang, and N. Xi, "A hybrid deep architecture for robotic grasp detection," in *Proceedings of IEEE International Conference on Robotics and Automation*, 2017, pp. 1609–1614.
- [24] E. Johns, S. Leutenegger, and A. J. Davison, "Deep learning a grasp function for grasping under gripper pose uncertainty," in *Proceedings of the IEEE International Conference on Intelligent Robotic and Systems*, 2016, pp. 4461–4468.
- [25] J. Mahler, J. Liang, S. Niyaz, M. Laskey, R. Doan, X. Liu, J. Aparicio-Ojea, and K. Goldberg, "Dex-Net 2.0: Deep learning to plan robust grasps with synthetic point clouds and analytic grasp metrics," in *Robotics: Science and Systems*, 2017.
- [26] Q. Lu, K. Chenna, B. Sundaralingam, and T. Hermans, "Planning multi-fingered grasps as probabilistic inference in a learned deep network," in *International Symposium on Robotics Research*, 2017.
- [27] L. Pinto and A. Gupta, "Supersizing self-supervision: Learning to grasp from 50k tries and 700 robot hours," in *Proceedings of IEEE International Conference on Robotics and Automation*, 2016, pp. 3406–3413.
- [28] S. Levine, P. Pastor, A. Krizhevsky, and D. Quillen, "Learning hand-eye coordination for robotic grasping with large-scale data collection," in *International Symposium on Experimental Robotics*, 2016, pp. 173–184.
- [29] S. Ren, K. He, R. Girshick, and J. Sun, "Faster R-CNN: Towards real-time object detection with region proposal networks," in *Advances in Neural Information Processing Systems*, 2015, pp. 91–99.
- [30] R. Girshick, J. Donahue, T. Darrell, and J. Malik, "Rich feature hierarchies for accurate object detection and semantic segmentation," in *Proceedings of the IEEE conference on computer vision and pattern recognition*, 2014, pp. 580–587.
- [31] R. Girshick, "Fast R-CNN," in *IEEE International Conference on Computer Vision*, 2015, pp. 1440–1448.
- [32] T.-Y. Lin, M. Maire, S. Belongie, J. Hays, P. Perona, D. Ramanan, P. Dollár, and C. L. Zitnick, "Microsoft coco: Common objects in context," in *European Conference on Computer Vision*, 2014, pp. 740–755.